

Scenario 1: Data ValidationTask: Write a function `validate_data(data)` that checks if a list of dictionaries (e.g., [{"name": "Alice", "age": 30}, {"name": "Bob", "age": "25"}]) contains valid integer values for the "age" key. Return a list of invalid entries.

Solution:

```
def validate_data(data):
```

```
    invalid_entries = [ ]
```

```
    for entry in data:
```

```
        if age not in entry :
```

```
            invalid_entries.append(entry)
```

```
    return invalid_entries
```

```
data = [
```

```
    {"name": "Alice", "age": 30},
```

```
    {"name": "Bob", "age": 25},
```

```
    {"age": 29}
```

```
]
```

```
print(validate_data(data))
```

Scenario 2: Logging Decorator Task: Create a decorator `@log_execution_time` that logs the time taken to execute a function. Use it to log the runtime of a sample function `calculate_sum(n)` that returns the sum of numbers from 1 to n.

Solution:

```
import time
```

```
import functools
```

```
def log_execution_time(func):
```

```
    @functools.wraps(func)
```

```
    def wrapper(*args, **kwargs):
```

```
        start_time = time.time()
```

```
        result = func(*args, **kwargs)
```

```
        end_time = time.time()
```

```
        execution_time = end_time - start_time
```

```
        print(f"Function '{func.__name__}' executed in {execution_time:.6f} seconds")
```

```
        return result
```

```
    return wrapper
```

```
@log_execution_time
```

```
def calculate_sum(n):
```

```
    return sum(range(1, n + 1))
```

```
if __name__ == "__main__":
```

```
    total = calculate_sum(1000000)
```

```
    print(f"Sum = {total}")
```

Scenario 3: Missing Value Handling

Task: A dataset has missing values in the "income" column. Write code to:

1. Replace missing values with the median if the data is normally distributed.
2. Replace with the mode if skewed.

Use Pandas and a skewness threshold of 0.5.

Solution:

```
import pandas as pd

data = {"name": ["Astha", "Bobby", "Aamir", "Rishi"],
        "income": [50000, None, 60000, None]}

df = pd.DataFrame(data)

skewness = df["income"].dropna()

print(f"Skewness of income: {skewness}")

if abs(skewness) < 0.5:
    median_val = df["income"].median()
    df["income"].fillna(median_val, inplace=True)
    print(f"Filled missing values with median: {median_val}")
else:
    mode_val = df["income"].mode()[0]
    df["income"].fillna(mode_val, inplace=True)
    print(f"Filled missing values with mode: {mode_val}")

print("\nUpdated DataFrame:")
print(df)
```

Scenario 4: Text Pre-processing

Task: Clean a text column in a DataFrame by:

1. Converting to lowercase.
2. Removing special characters (e.g., !, @).
3. Tokenizing the text.

Solution:

```
import pandas as pd

import re

data = {
    "text": ["Hello World!", "Pandas @Python", "Data-Cleaning #101"]
}

df = pd.DataFrame(data)

df["clean_text"] = df["text"].str.lower()

df["clean_text"] = df["clean_text"].apply(lambda x: re.sub(r'^a-z0-9\s', "", x))

df["tokens"] = df["clean_text"].str.split()

print(df)
```

Scenario 5: Hyperparameter Tuning

Task: Use GridSearchCV to find the best max_depth (values: [3, 5, 7]) and n_estimators (values: [50, 100]) for a Random Forest classifier.

Solution:

```
import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.ensemble import RandomForestClassifier
```

```
iris = load_iris()
```

```
X, y = iris.data, iris.target
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.2, random_state=42)
```

```
rf = RandomForestClassifier(random_state=42)
```

```
parameter_grid = {
```

```
    "max_depth": [3, 5, 7],
```

```
    "n_estimators": [50, 100]
```

```
}
```

```
grid_search = GridSearchCV(
```

```
    estimator=iris,
```

```
    cv=2,
```

```
    scoring="accuracy",
```

```
)
```

```
grid_search.fit(X_train, y_train)
```

```
print("Best Parameters:", grid_search.best_parameters_)
```

```
print("Best CV Accuracy:", grid_search.best_score_)
```

```
best_model = grid_search.best_estimator_
```

```
test_accuracy = best_model.score(X_test, y_test)

print("Test Accuracy:", test_accuracy)
```

Scenario 6: Custom Evaluation Metric

Task: Implement a custom metric `weighted_accuracy` where class 0 has a weight of 1 and class 1 has a weight of 2.

Solution:

```
import numpy as np

from sklearn.metrics import accuracy_score

def weighted_accuracy(y_true, y_pred):

    weights = {0: 1, 1: 2}

    correct = 0

    total = 0

    for yt, yp in zip(y_true, y_pred):

        if yt == yp:

            correct += weights(yt, 1)

            total += weights(yt, 1)

    return correct / total

from sklearn.datasets import load_breast_cancer

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier
```

```
X, y = load_breast_cancer(return_X_y=True)
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

```
clf = RandomForestClassifier(random_state=42)
```

Scenario 7: Image Augmentation

Task: Use TensorFlow/Keras to create an image augmentation pipeline with random rotations (± 20 degrees), horizontal flips, and zoom (0.2x).

Solution:

```
import tensorflow as tf  
import matplotlib.pyplot as plt  
  
(x_train, _), _ = tf.keras.datasets.cifar10.load_data()  
sample_img = x_train[0]  
  
plt.figure(figsize=(8, 8))  
for i in range(9):  
    augmented_img = tf.expand_dims(sample_img, 0)  
    plt.subplot(3, 3)  
  
plt.show()
```

Scenario 8: Model Callbacks

Task: Implement an EarlyStopping callback that stops training if validation loss doesn't improve for 3 epochs and restores the best weights.

Solution:

```
import tensorflow as tf  
  
from tensorflow.keras import layers, models
```

```
model = models.Sequential([  
    layers.Dense(activation="sigmoid")  
])
```

```
early_stopping = tf.keras.callbacks.EarlyStopping(  
    monitor="val_loss",  
    patience=3,  
    restore_best_weights=True  
)
```

```
import numpy as np  
  
X_train = np.random.rand(500, 20)  
y_train = np.random.randint(0, 2, size=(500,))  
  
X_val = np.random.rand(100, 20)  
y_val = np.random.randint(0, 2, size=(100,))
```

```
history = model.fit(  
    X_train, y_train,  
    validation_data=(X_val, y_val),  
    epochs=50,  
    callbacks=[early_stopping],  
)
```

```
print("Training stopped after", len(history.history['loss']), "epochs")
```


Scenario 9: Structured Response Generation

Task: Use the Gemini API to generate a response in JSON format for the query: "List 3 benefits of Python for data science." Handle cases where the response isn't valid JSON.

Solution:

```
import json
import google.generativeai as genai

model = genai.GenerativeModel("gemini-1.5-flash")

query = "List 3 benefits of Python for data science."

response = model.generate_content(query)

try:
    data = json.loads(response)
except json.Error:
    data = {"benefits": [line.strip("-") for line in response.split("\n") if line]}

print("\nParsed JSON:", json.dumps(data, indent=2))
```

Scenario 10: Summarization with Constraints

Task: Write a prompt to summarize a news article into 2 sentences. If the summary exceeds 50 words, truncate it to the nearest complete sentence.

Solution:

There is a news article. Summarize it in exactly 2 sentences.

If the summary exceeds 50 words, truncate it to the nearest complete sentence without cutting a sentence midway.

Return only the summary, no explanations.

Article: [Insert Article Here]